

# FLAVOURBENCH: What Language Models Know and What Epicure Adds

Josef Chen  
Independent Researcher

August 2026

## Abstract

Culinary benchmarks usually rely on static question banks or model judges. FLAVOURBENCH uses a versioned EPICURE runtime instead. Four deterministic EPICURE operations produce 32 multiple-choice tasks and their answer keys before any evaluated model is called. No language model writes or judges the test. For this benchmark, ground truth means agreement with that pinned runtime.

Each of 20 model endpoints receives every task twice. In the *Model only* condition, exact accuracy is the FlavourBench Score. In the *Model + Epicure* condition, the model may make one read-only call to the named operation. The difference is Epicure Gain. The release contains 640 matched pairs and 1280 scored responses. GPT-5.6 Sol Pro leads with a FlavourBench Score of 62.5. Across the panel, Epicure changes accuracy from 43.9% to 93.9%, a gain of +50.0 percentage points. The assisted condition measures whether a model can use a specified operation, not whether it can discover the right tool.

A single content-addressed release contains the tasks, answer keys, routes, responses, and bounded Epicure traces. It reconstructs every score, table, and figure offline. FLAVOURBENCH therefore separates an operational model score from the measured benefit of a domain system, without a human or model judge.

## 1 Introduction

Benchmarks are easiest to trust when their ground truth is external to the models being tested. Code can be executed, database states can be inspected, and games can be scored. Culinary reasoning rarely has that luxury. Ingredient substitutions, pairings, technique choices, and cuisine associations are commonly assessed by another language model or by a small set of human preferences. The evaluator then becomes part of the construct.

FLAVOURBENCH starts from EPICURE, a read-only culinary runtime with 1,790 ingredients represented in a 300-dimensional space [16, 17]. Four runtime operations are converted into four-choice questions. The compiler records each answer before it contacts an evaluated model.

The design asks two separate questions. The FlavourBench Score measures how often a model endpoint returns EPICURE’s answer without access to the runtime. Epicure Gain measures the change when the same endpoint receives the same task with one read-only EPICURE call. Every model receives every task in both conditions. The score determines the leaderboard; Epicure Gain does not enter the ranking.

Our release evaluates 20 current endpoints, including GPT-5.6 Sol, Terra, and Luna; Claude Fable, Opus, and Sonnet 5; Gemini 3.1 Pro and 3.6 Flash; Grok 4.5; Kimi K3; Qwen 3.8 Max; GLM 5.2; two DeepSeek V4 variants; MiniMax M3; Nemotron 3 Ultra; Mistral Large 2512; Tencent HY 3; and Cohere Command A Plus and Command A Reasoning. Kimi and both Cohere models use direct provider routes. Every other endpoint uses an exact, no-fallback OpenRouter route selected before the

run.

The paper makes three contributions:

1. an automated culinary benchmark whose questions and answer keys come from a versioned executable system rather than a model judge;
2. a common-task frontier evaluation with 640 matched pairs, exact missing-response penalties, and a public 20-row leaderboard;
3. a matched Epicure intervention that reports panel gain and task-family interactions without changing the ranking score.

EPICURE has two roles. It fixes the answer key, then becomes the intervention in the second arm. This keeps the leaderboard and the Epicure comparison connected without combining them into an arbitrary composite score.

## 2 Epicure reference answers

### 2.1 A versioned reference

For FLAVOURBENCH, ground truth is the output of a specific version of EPICURE. It is not a claim about universal culinary preference. The release binds the application, evidence bundle, release identifier, ingredient count, and embedding dimensionality by hash. A different runtime produces a different task-set hash.

This definition makes every answer independently recomputable. It also gives the score a direct interpretation: accuracy is agreement with the published culinary runtime under the published task construction. No second model decides which response sounds better.




Figure 1: The FLAVOURBENCH pipeline. EPICURE fixes each answer before evaluation. The Model only arm produces the ranked FlavourBench Score. The matched Model + Epicure arm produces Epicure Gain. Both are reconstructed from the same public release.

## 2.2 Four balanced task families

The compiler creates eight tasks in each of four families. Seeds and ingredient pairs are fixed before evaluation. For every task, EPICURE is called once to obtain the reference result. The correct answer is extracted programmatically, three distractors are constructed without model generation, and the four options are shuffled with a task-ID-derived seed.

The task families deliberately mix retrieval and computation. Neighbour and cuisine questions test whether a model knows relationships encoded by the graph. Pairing questions require an exact numeric result that is difficult to guess. Axis questions test a comparative claim rather than a memorised scalar. Together they prevent the leaderboard from reducing to one operation or one kind of culinary trivia.

## 2.3 Prompt and answer contract

Each prompt names the relevant operation and gives four choices. Models may explain briefly, but the last answer line must be `FINAL_CHOICE: X`. The parser accepts A, B, C, or D on that marker line. A missing marker, abnormal provider finish, timeout, or absent response scores zero. The rule fits in the standalone verifier shipped with the release.

The complete task set contains no hidden human label and no model-authored item. The benchmark is therefore suitable for automatic reruns, CI-style regression tests, and training loops where thousands of candidate responses must be scored consistently.

# 3 Evaluation design

## 3.1 One score and one matched intervention

For model  $m$  and task  $t$ , let  $M_{mt}$  equal one when the Model only response finishes normally and matches the EPICURE answer. Let  $E_{mt}$  be the same indicator for Model + Epicure. The prompt, route, temperature, top- $p$ , seed request, output limit, and answer parser are identical. Only access to EPICURE changes.

Each Model + Epicure response may make one read-only call to the operation named in the prompt. The runner records the operation, arguments, latency, error flag, and a hash of the returned payload. The returned data is bounded before it enters model context. Automatic provider fallback is disabled, so a route failure remains a failure for that route.

The single ranking metric is the FlavourBench Score:

$$F_m = 100 \frac{1}{32} \sum_t M_{mt}. \quad (1)$$

Every family contains eight tasks, so ordinary accuracy already gives each family equal weight. Missing or invalid responses remain in the denominator as incorrect. The public release sorts by  $F_m$ , then uses Model + Epicure accuracy, Model + Epicure completion, and model ID as fixed tie-breakers.

Epicure Gain is the matched difference

$$G_m = 100 \frac{1}{32} \sum_t (E_{mt} - M_{mt}), \quad (2)$$

reported in percentage points.  $G_m$  is not part of the FlavourBench Score. The release also keeps the four paired outcomes for every task: both correct, Model only correct, Model + Epicure correct, or neither correct.

## 3.2 Completion and uncertainty

Completion is the fraction of assigned responses that end normally with a parseable final marker. The score still counts missing responses as incorrect. Reporting completion beside the score keeps endpoint failure separate from a wrong culinary answer. A zero score with zero completions is not evidence of zero model capability. For each 32-task accuracy we report a 95% Wilson interval.

The benchmark is designed for visible differences, not hundredths of a point. One answer changes the score by 3.125 points. Family breakdowns and the full pair matrix are published so tied or nearby aggregate scores can be inspected rather than over-interpreted.

## 3.3 Twenty-model frontier panel

The panel was fixed before the final calls. It includes three GPT-5.6 tiers rather than treating one deployment as the

| Family       | Epicure operation       | Exact target              | Example query                  | Tasks |
|--------------|-------------------------|---------------------------|--------------------------------|-------|
| Substitution | <b>neighbors</b>        | first ranked neighbour    | top neighbour for tomato       | 8     |
| Composition  | <b>pairing_score</b>    | four-decimal affinity     | tomato + basil                 | 8     |
| Cookability  | <b>compare_on_axis</b>  | larger flavour projection | coffee vs. chocolate on bitter | 8     |
| Evidence     | <b>cultural_profile</b> | highest cuisine cosine    | cuisine direction for kimchi   | 8     |

Table 1: The task set is balanced by construction. Each family maps one deterministic EPICURE operation to a four-choice answer. Chance accuracy is 25%.

entire OpenAI family; three Claude 5 tiers; two Gemini tiers; two DeepSeek V4 routes; both current Cohere Command A routes; and a broad set of current independent and open-weight families. Kimi K3 runs through Kimi’s Anthropic-compatible endpoint. Command A Plus and Command A Reasoning run through Cohere V2 Chat. Qwen 3.8 Max and GLM 5.2 use explicit OpenRouter routes in the automated batch [9, 12, 3, 5, 4, 13]. The exact 20-row route table appears in Appendix A.

## 4 Results

### 4.1 The FlavourBench Score leaderboard

GPT-5.6 Sol Pro leads with 20/32 correct and a FlavourBench Score of 62.5. Table 2 gives the full release ordering and completion rates. Figure 2 adds intervals to show which gaps the 32-task panel can resolve. Models with no Model only completion are marked DNF rather than interpreted as the least capable systems.

The answer key exists before the model call, every endpoint receives the same questions, and no judge is sampled after the fact. Missing outputs count as incorrect in the operational score but remain visible in the completion columns. Mistral Large 2512 and Nemotron 3 Ultra returned no Model only completion in this run, so their zeroes are transport outcomes rather than capability estimates. Nemotron also returned no Model + Epicure completion.

### 4.2 What Epicure changes

Across all 640 matched pairs, panel accuracy is 43.9% for Model only and 93.9% for Model + Epicure. Epicure Gain is therefore +50.0 points. Every route that produced at least one assisted response gains points under the operational scoring rule. The largest numerical change is Mistral Large 2512 at +93.8 points, but that route had no Model only completion. Its change combines condition-specific availability with answer quality and should not be read as a pure knowledge gain.

The assisted prompt names the relevant operation. This is a scaffolded test of calling that operation, reading its output, and preserving the answer contract. It is not a test of open-ended tool discovery. That narrower intervention is useful here because it isolates the benefit of EPICURE from the separate problem of choosing among many tools.

### 4.3 Families explain aggregate ties

The four task families have different demands. Pairing questions require an exact four-decimal value; substitution and cultural questions can sometimes be answered from broad food knowledge; axis comparisons require a directional decision in EPICURE’s latent space. Figure 4 shows that similar total scores can arise from different profiles.

### 4.4 Every model and task pair is inspectable

Aggregate bars can hide a benchmark bug or a single pathological task. Figure 5 therefore publishes all 640 paired outcomes. Blue cells are stable knowledge: both conditions are correct. Gold cells are Epicure wins: Model only is wrong and Model + Epicure is correct. Red cells show the reverse outcome. Grey cells show questions neither condition solved.

### 4.5 Cohere, Kimi, Qwen, and route fidelity

Both Cohere models are full members of the public panel. Command A Plus ranks 16 with a FlavourBench Score of 40.6 and an Epicure Gain of +59.4 points. Command A Reasoning ranks 12 with a score of 46.9 and a gain of +50.0 points. Kimi K3 ranks 5 through the direct Kimi endpoint. Qwen 3.8 Max ranks 11 through an exact OpenRouter route, and GLM 5.2 ranks 15. No route changes during generation.

The two direct Cohere routes each produced all 64 assigned arms. Their returned identifiers are `command-a-plus-05-2026` and `command-a-reasoning-08-2025`, respectively; the direct Kimi route likewise produced all 64 arms under returned identifier `k3`. The offline verifier checks these provider, identity, and arm-count contracts before accepting the release.

This matters because a model name is not a transport contract. The release stores the requested model, returned model, provider, route, generation ID, latency, token counts, and response hash. Direct and routed endpoints can therefore be analysed separately without pretending they are interchangeable deployments.




Figure 2: FlavourBench Score against EPICURE answer keys. Points are exact accuracy over the common 32-task panel; bars are 95% Wilson intervals. Crosses mark endpoints with no Model only completion. The dashed line is four-choice chance.

#### 4.6 Score, response time, and Epicure Gain

Figure 6 plots FlavourBench Score against median response time. Marker colour encodes Epicure Gain. A fast model can score poorly, a high-scoring model can gain little from EPICURE, and a lower-ranked model can improve sharply when assisted. The release keeps these dimensions separate.

### 5 Real prompts, answers, and tool calls

A benchmark paper should show what was actually asked. The four cases below are selected mechanically, one per family. Where possible, the builder chooses a model and task that is wrong in Model only and correct in Model + Epicure. Each case includes the full prompt, both final

answers, the actual operation and arguments, and the exact answer key.

**Case 1: Substitution with GPT-5.6 Sol Pro. Prompt.** FlavourBench Epicure-native exact task. Use the read-only culinary evidence tools if they are available; otherwise answer from your own knowledge.

According to Epicure’s ‘neighbors’ result for ‘tomato’ with top\_k=4, which ingredient is ranked first?

Choices: A. oregano B. bell\_pepper C. olive\_oil D. onion

In your final answer, explain briefly if useful and finish on its own line with exactly ‘FINAL\_CHOICE: X’, replacing X with A, B, C, or D.

**Model only.** FINAL\_CHOICE: D **Model + Epicure.** FINAL\_CHOICE: B

**Epicure call.** neighbors({"ingredient": "tomato", "top\_k": 4}). The answer key is

| Rank | Model               | FB Score | +Epicure | Gain  | Model % done | Epicure % done |
|------|---------------------|----------|----------|-------|--------------|----------------|
| 1    | GPT-5.6 Sol Pro     | 62.5     | 100.0    | +37.5 | 96.9         | 100.0          |
| 2    | Claude Sonnet 5     | 59.4     | 100.0    | +40.6 | 100.0        | 100.0          |
| 3    | DeepSeek V4 Pro     | 59.4     | 93.8     | +34.4 | 100.0        | 93.8           |
| 4    | Gemini 3.6 Flash    | 56.2     | 100.0    | +43.8 | 100.0        | 100.0          |
| 5    | Kimi K3             | 56.2     | 100.0    | +43.8 | 100.0        | 100.0          |
| 6    | Tencent HY 3        | 56.2     | 100.0    | +43.8 | 100.0        | 100.0          |
| 7    | Claude Fable 5      | 53.1     | 100.0    | +46.9 | 93.8         | 100.0          |
| 8    | Claude Opus 5       | 50.0     | 100.0    | +50.0 | 100.0        | 100.0          |
| 9    | GPT-5.6 Luna Pro    | 46.9     | 100.0    | +53.1 | 100.0        | 100.0          |
| 10   | GPT-5.6 Terra Pro   | 46.9     | 100.0    | +53.1 | 100.0        | 100.0          |
| 11   | Qwen 3.8 Max        | 46.9     | 100.0    | +53.1 | 87.5         | 100.0          |
| 12   | Command A Reasoning | 46.9     | 96.9     | +50.0 | 93.8         | 100.0          |
| 13   | Gemini 3.1 Pro      | 46.9     | 96.9     | +50.0 | 93.8         | 96.9           |
| 14   | MiniMax M3          | 43.8     | 100.0    | +56.2 | 96.9         | 100.0          |
| 15   | GLM 5.2             | 43.8     | 96.9     | +53.1 | 100.0        | 96.9           |
| 16   | Command A Plus      | 40.6     | 100.0    | +59.4 | 100.0        | 100.0          |
| 17   | DeepSeek V4 Flash   | 40.6     | 100.0    | +59.4 | 75.0         | 100.0          |
| 18   | Grok 4.5            | 21.9     | 100.0    | +78.1 | 31.2         | 100.0          |
| DNF  | Mistral Large 2512  | 0.0      | 93.8     | +93.8 | 0.0          | 100.0          |
| DNF  | Nemotron 3 Ultra    | 0.0      | 0.0      | +0.0  | 0.0          | 0.0            |

Table 2: The public leaderboard. “FB Score” is Model only exact-choice accuracy and determines rank. “+Epicure” is accuracy for the same tasks with EPICURE available. “Gain” is their difference in percentage points. The final columns report completion. DNF means that the Model only endpoint returned no parseable completion.

| Family       | Epicure operation | Tasks | Model only | Model + Epicure |
|--------------|-------------------|-------|------------|-----------------|
| Substitution | neighbors         | 8     | 18.1       | 94.4            |
| Composition  | pairing_score     | 8     | 40.6       | 94.4            |
| Cookability  | compare_on_axis   | 8     | 62.5       | 92.5            |
| Evidence     | cultural_profile  | 8     | 54.4       | 94.4            |

Table 3: Panel-average performance by family and the corresponding reference operation.

FINAL\_CHOICE: B.

**Epicure result.** {"ingredient": "tomato", "neighbors": [{"name": "bell\_pepper", "rank": 1, "sim": 0.4799}, {"name": "olive\_oil", "rank": 2, "sim": 0.4786}, {"name": "oregano", "rank": 3, "sim": 0.4646}, {"name": "onion", "rank": 4, "sim": 0.4567}]}

**Case 2: Composition with DeepSeek V4 Pro. Prompt.** FlavourBench Epicure-native exact task. Use the read-only culinary evidence tools if they are available; otherwise answer from your own knowledge.

What exact four-decimal ‘pairing\_score’ does Epicure return for ‘tomato’ and ‘basil’?

Choices: A. 0.2707 B. 0.2407 C. 0.3207 D. 0.1907

In your final answer, explain briefly if useful and finish on its own line with exactly ‘FINAL\_CHOICE: X’, replacing X with A, B, C, or D.

**Model only.** FINAL\_CHOICE: B **Model + Epicure.** FINAL\_CHOICE: A

**Epicure call.** pairing\_score({"ingredient\_a": "tomato", "ingredient\_b": "basil"}). The answer key is FINAL\_CHOICE: A.

**Epicure result.** {"all\_pairs\_range": {"median": 0.092, "p10": 0.0197, "p25": 0.0521, "p75": 0.1392, "p90": 0.1899}, "pairing\_score": 0.2707, "percentile\_label": "very high (>=p90)", "resolved\_a": "tomato", "resolved\_b": "basil"}

**Case 3: Cookability with GPT-5.6 Sol Pro. Prompt.** FlavourBench Epicure-native exact task. Use the read-only culinary evidence tools if they are available; otherwise answer from your own knowledge.

Which statement matches Epicure’s ‘compare\_on\_axis’ result for ‘coffee’ versus ‘chocolate’ on ‘cf\_bitter’?

Choices: A. coffee has the higher projection B. the projections are equal to four decimals C. Epicure cannot resolve one or both ingredients D. chocolate has the higher projection

In your final answer, explain briefly if useful and finish on its own line with exactly ‘FINAL\_CHOICE: X’, replacing X with A, B, C, or D.

**Model only.** FINAL\_CHOICE: A **Model + Epicure.** FINAL\_CHOICE: D

**Epicure call.** compare\_on\_axis({"axis": "cf\_bitter", "ingredient\_a": "coffee", "ingredient\_b": "chocolate"}). The answer key is FINAL\_CHOICE: D.

**Epicure result.** {"axis": "cf\_bitter", "axis\_range": {"p10": -0.0877, "p90": 0.1429}, "delta": 0.0835, "label\_a": "moderate", "label\_b": "high", "projection\_a": 0.0076, "projection\_b": 0.0911, "resolved\_a": "coffee", "resolved\_b": "chocolate"}

**Case 4: Evidence with Grok 4.5. Prompt.** FlavourBench Epicure-native exact task. Use the read-only culinary evidence tools if they are available; otherwise answer from your own knowledge.

Which cuisine direction has the highest cosine score in Epicure’s ‘cultural\_profile’ result for ‘miso’?

Choices: A. Japanese B. East\_Asian C. South-east\_Asian D. South\_Asian

In your final answer, explain briefly if useful and finish on its own line with exactly ‘FINAL\_CHOICE: X’, replacing X with A, B, C, or D.

**Model only.** {"name": "cultural\_profile",




Figure 3: Model only and Model + Epicure accuracy over the same 32 tasks. Each line is a matched comparison within one endpoint, not a comparison between different task samples.

"arguments": {"ingredient": "miso"}} **Model**  
+ **Epicure**. FINAL\_CHOICE: A

**Epicure call.** cultural\_profile({"ingredient":  
"miso"}). The answer key is FINAL\_CHOICE: A.

**Epicure result.** {"cuisines": {"East\_Asian":  
{"range": {"p10": -0.1649, "p90": 0.3039}, "score":  
0.1625}, "Eastern\_European": {"range": {"p10": -0.1466,  
"p90": 0.1034}, "score": -0.1494}, "Japanese":  
{"range": {"p10": -0.1359, "p90": 0.1896}, "score":  
0.3227}, "Latin\_American": {"range": {"p10": -0.1672,  
"p90": 0.1281}, "score": -0.1538}, "Mediterranean":  
{"range": {"p10": -0.2323, "p90": 0.206}, "score":  
-0.1198}, "South\_Asian": {"range": {"p10": -0.0808,  
"p90": 0.071}, "score": -0.0585}, "Southeast\_Asian":  
{"range": {"p10": -0.1625, "p90": 0.193}, "score":  
0.0426}, "Western\_Atlantic": {"range": {"p10": -0.2304,  
"p90": 0.183}, "score": -0.1546}, "note": "Cosine  
similarity to each cuisine direction. Use p10/p90 to  
judge whether a score is typical or extreme.", "resolved":

"miso"}

These examples show exactly what the two measure-  
ments contain. Model only tests knowledge or inference  
about the hidden graph. Model + Epicure names the  
operation, so it tests whether the model can form valid  
arguments, read the returned structure, and preserve the  
answer contract. A correct choice with the wrong opera-  
tion still earns answer credit, but the release separately  
reports whether the reference operation was matched.  
Across the run there are 606 projected EPICURE calls  
and 606 model and task pairs matching the reference  
operation.

## 6 Using the score as a training sig- nal

The same environment exposes four binary indicators for  
a generated trajectory  $\tau$ :  $r_{\text{answer}}$  for a correct answer,

$r_{\text{format}}$  for a parseable final marker,  $r_{\text{call}}$  for a successful Epicure call, and  $r_{\text{align}}$  for a call matching the reference operation. The public leaderboard uses  $r_{\text{answer}}$  in Model only. A future training objective could use

$$r(\tau) = r_{\text{answer}}r_{\text{format}} + \lambda r_{\text{answer}}r_{\text{call}} + \mu r_{\text{align}}, \quad (3)$$

where each component is a binary indicator and  $\lambda, \mu$  set the value of successful and reference-matching calls. Requiring a correct answer prevents irrelevant calls from earning the main reward. The interface could support rejection sampling, preference optimisation, or reinforcement learning on new ingredient seeds. This paper does not train a checkpoint; such a study should keep the public leaderboard as a final evaluation split.

## 7 Reproducibility and release

The public release is one JSON document with hash 197b84d75cca.... It contains the 20 route records, 32 full prompts, answer keys and reference EPICURE results, 1,280 observation rows, final answer markers, response times, and bounded Epicure-trace projections. 1195 provider responses were observed; 85 missing arms remain explicit zero-scored rows.

The leaderboard can be reconstructed without a network connection, provider key, GPU, or running EPICURE service:

```
python3 -I reproduce_epicure_native.py \
--release generated/epicure-native/\
epicure-native-release.json
```

The verifier rejects duplicate JSON keys, non-finite numbers, a changed content address, an incomplete  $20 \times 32 \times 2$  grid, or any leaderboard value that differs from recomputation. The paper tables and figures are generated from the same release, so a changed score cannot leave a stale chart behind.

Reproducing the original hosted calls is a different operation from replaying the result. The release records exact routes and request policies, but provider deployments can change. A fresh frontier run should generate a new manifest and release rather than overwrite this snapshot.

## 8 Discussion

**What the score means.** The FlavourBench Score answers a narrow operational question: how often did this deployed endpoint return the answer fixed by the published EPICURE runtime? It includes endpoint availability because missing responses score zero. The completion columns reveal that component. The score is not a universal measure of cooking skill, but it is exact, reproducible, and independent of a model judge.

**Do we need humans?** Not for this leaderboard. Humans are useful when the target is preference, creativity, safety, or the quality of an open-ended recipe. They are unnecessary when the question is whether a model’s choice matches a deterministic runtime result. FLAVOURBENCH uses that distinction to automate the core ranking while leaving subjective culinary evaluation as a separate track.

**Why run Model + Epicure?** The leaderboard says which endpoint best predicts EPICURE. The matched intervention asks how much the same endpoint changes when given EPICURE directly. That result is operationally useful: a strong model may gain little, while a lower-scoring model may become competitive when grounded. The dumbbell and pair matrix show this task by task.

**Scaling the benchmark.** The task compiler can expand the seed lists while preserving the four operations. The next release should reserve unseen ingredients and pairs, remove near-duplicate questions, and measure stability across runtime versions. The current 32-task panel is a frontier snapshot and a reproducible template for a larger hidden set.

## 9 Related work

HELM argues for explicit scenarios and metrics, while recent benchmark audits emphasise validity, contamination, and transparent reporting [10, 19]. Chatbot Arena and WildBench evaluate broad answer quality through preferences or checklists [1, 11]. FLAVOURBENCH instead targets an executable domain model and needs no judge for its primary score.

ToolLLM, BFCL, and  $\tau$ -bench make tool invocation or environment state observable [15, 14, 24]. FLAVOURBENCH adds a matched Model only condition, so EPICURE becomes an intervention rather than a requirement shared by every response. LiveBench, LiveCodeBench, GPQA, MMLU-Pro, and SWE-bench illustrate complementary strategies for reducing contamination or grounding answers in difficult tasks and executable outcomes [22, 7, 18, 21, 8].

Culinary datasets such as FlavorDB and Recipe1M organise ingredient and recipe evidence [6, 20]. CookingSense and planning-oriented culinary work test commonsense and procedural reasoning [2, 23]. EPICURE occupies a different layer: it exposes a callable structured representation whose outputs can be compiled into exact evaluation targets.

## 10 Limitations

The 32-task panel resolves large score differences better than close ones: one answer moves the score by 3.125 points. The English ingredient seeds cover four EPICURE

| Benchmark     | Scope              | Per-answer target           | Deterministic score | Matched Epicure intervention |
|---------------|--------------------|-----------------------------|---------------------|------------------------------|
| Chatbot Arena | general chat       | pairwise preference         | No                  | No                           |
| WildBench     | general chat       | task checklist              | No                  | No                           |
| BFCL          | function calling   | structured call             | Yes                 | No                           |
| $\tau$ -bench | interactive agents | exact goal state            | Yes                 | No                           |
| FLAVOURBENCH  | culinary reasoning | versioned EPICURE operation | Yes                 | Yes                          |

Table 4: FLAVOURBENCH complements broad preference and tool-use benchmarks. Its distinctive unit is the same endpoint and task measured as Model only and Model + Epicure.

operations, and hosted routes are a dated snapshot. Public prompts may eventually be memorised. The assisted condition names the operation, so it measures execution with a specified tool rather than open-ended tool selection. Finally, agreement with EPICURE inherits the runtime’s representation choices; it is not a food-safety or cultural-authenticity certification. Larger hidden panels, additional operations, and repeated route snapshots are the immediate next tests.

## 11 Conclusion

FLAVOURBENCH turns a versioned culinary runtime into a public automated benchmark. GPT-5.6 Sol Pro leads the 20-model panel with a FlavourBench Score of 62.5. Giving the same endpoints a specified EPICURE operation raises panel accuracy by +50.0 points. The score, gain, completion rates, prompts, responses, and traces all reconstruct from the released 1,280-arm grid. The result is a simple benchmark with a separate, measured intervention: what models know, and what EPICURE adds.




Figure 4: Model only and Model + Epicure accuracy by task family. Each value is a percentage over eight tasks. Similar FlavourBench Scores can come from different family profiles.




Figure 5: All 640 matched outcomes, ordered by FlavourBench Score and grouped by task family. The matrix shows which exact tasks produce Epicure Gain and where the assisted response regresses.




Figure 6: FlavourBench Score, response time, and Epicure Gain. Response time uses a log scale. Point size increases with the magnitude of gain; colour indicates its direction. An  $\times$  marks a route with no observed response time.

## A Exact execution routes

| Slot | Model               | Execution route                          |
|------|---------------------|------------------------------------------|
| 1    | GPT-5.6 Sol Pro     | openrouter: openai/flex                  |
| 6    | Claude Sonnet 5     | openrouter: amazon-bedrock/claude-on-aws |
| 13   | DeepSeek V4 Pro     | openrouter: baidu/fp8                    |
| 8    | Gemini 3.6 Flash    | openrouter: google-vertex/global/flex    |
| 10   | Kimi K3             | kimi_direct: kimi-code-direct            |
| 18   | Tencent HY 3        | openrouter: baidu/fp8                    |
| 4    | Claude Fable 5      | openrouter: amazon-bedrock               |
| 5    | Claude Opus 5       | openrouter: amazon-bedrock/claude-on-aws |
| 3    | GPT-5.6 Luna Pro    | openrouter: openai/flex                  |
| 2    | GPT-5.6 Terra Pro   | openrouter: openai/flex                  |
| 11   | Qwen 3.8 Max        | openrouter: alibaba                      |
| 20   | Command A Reasoning | cohere_direct: cohere-direct             |
| 7    | Gemini 3.1 Pro      | openrouter: google-vertex/global/flex    |
| 15   | MiniMax M3          | openrouter: together                     |
| 12   | GLM 5.2             | openrouter: coreweave/fp4                |
| 19   | Command A Plus      | cohere_direct: cohere-direct             |
| 14   | DeepSeek V4 Flash   | openrouter: cloudflare/fp8               |
| 9    | Grok 4.5            | openrouter: xai/zdr                      |
| 17   | Mistral Large 2512  | openrouter: mistral                      |
| 16   | Nemotron 3 Ultra    | openrouter: deepinfra/fp4                |

Table 5: All 20 model routes in the released leaderboard. Routes are selected before generation; automatic fallback is disabled.

## B Protocol details

All arms request temperature 0, top- $p$  1, and seed 20260810 where the provider supports it. The final response budget is 2,048 tokens. Model + Epicure has one call round, one call total, a 4,096-byte result bound, and a 4,096-byte cumulative result bound. At most two provider attempts are allowed, and an ambiguously delivered request is never replayed. The final marker is parsed from plain text for transport portability; direct providers may use native structured generation internally and normalize to the same marker.

## C Data fields

Each observation binds model ID, task ID, condition, source status, source and response hashes, returned model and provider, finish reason, response time, final answer, observed and expected choice, normal-parse flag, correctness, and a projected Epicure trace. Each trace event contains its name, structured arguments, error flag, response time, and result hash. The release does not require raw provider credentials or hidden service state.

## References

- [1] Wei-Lin Chiang, Lianmin Zheng, Ying Sheng, Anatasios Nikolas Angelopoulos, Tianle Li, Dacheng Li, Hao Zhang, Banghua Zhu, Michael I. Jordan, Joseph E. Gonzalez, and Ion Stoica. Chatbot Arena: An Open Platform for Evaluating LLMs by Human Preference. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 8359–8388. PMLR, 2024.
- [2] Donghee Choi, Mogan Gim, Donghyeon Park, Mujeen Sung, Hyunjae Kim, Jaewoo Kang, and Jihun Choi. CookingSense: A Culinary Knowledgebase with Multidisciplinary Assertions. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 3983–3996, Torino, Italia, 2024. ELRA and ICCL.
- [3] Cohere. Chat (V2) API Reference. <https://docs.cohere.com/v2/reference/chat>, 2026. Accessed 3 August 2026.
- [4] Cohere. Reasoning models. <https://docs.cohere.com/v2/docs/reasoning>, 2026. Accessed 3 August 2026.
- [5] Cohere. Tool use overview. <https://docs.cohere.com/v2/docs/tool-use-overview/>, 2026. Accessed 3 August 2026.
- [6] Neelansh Garg, Apuroop Sethupathy, Rudraksh Tuwani, et al. FlavorDB: A Database of Flavor Molecules. *Nucleic Acids Research*, 46(D1):D1210–D1216, 2018.
- [7] Naman Jain, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar-Lezama, Koushik Sen, and Ion Stoica. Live-CodeBench: Holistic and Contamination Free Evaluation of Large Language Models for Code. *arXiv preprint arXiv:2403.07974*, 2024.
- [8] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? In *The Twelfth International Conference on Learning Representations*, 2024.
- [9] Kimi Team. Kimi K3: Open Frontier Intelligence. *arXiv:2607.24653*, 2026.
- [10] Percy Liang, Rishi Bommasani, Tony Lee, et al. Holistic Evaluation of Language Models. *Transactions on Machine Learning Research*, 2023.
- [11] Bill Yuchen Lin, Yuntian Deng, Khyathi Chandu, Faeze Brahman, Abhilasha Ravichander, Valentina Pyatkin, Nouha Dziri, Ronan Le Bras, and Yejin Choi. WildBench: Benchmarking LLMs with Challenging Tasks from Real Users in the Wild. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [12] Moonshot AI. Kimi code: Providers and models. <https://moonshotai.github.io/kimi-code/en/configuration/providers.html>, 2026. Accessed 28 July 2026.
- [13] OpenRouter. Get Request and Usage Metadata for a Generation. <https://openrouter.ai/docs/api/reference/overview>, 2026. Accessed 2026-07-15.
- [14] Shishir G. Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. The Berkeley Function Calling Leaderboard (BFCL): From Tool Use to Agentic Evaluation of Large Language Models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 48371–48392. PMLR, 2025.
- [15] Yujia Qin, Shihao Liang, Yining Ye, et al. Tool-LLM: Facilitating Large Language Models to Master 16000+ Real-world APIs. In *The Twelfth International Conference on Learning Representations*, 2024.
- [16] Jakub Radzikowski and Josef Chen. Epicure: Multidimensional Flavor Structure in Food Ingredient Embeddings. *arXiv preprint arXiv:2604.22776*, 2026.
- [17] Jakub Radzikowski and Josef Chen. Epicure: Navigating the Emergent Geometry of Food Ingredient Embeddings. *arXiv preprint arXiv:2605.22391*, 2026.
- [18] David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R. Bowman. GPQA: A Graduate-Level Google-Proof Q&A Benchmark. In *First Conference on Language Modeling*, 2024.
- [19] Anka Reuel, Amelia Hardy, Chandler Smith, Max Lamparth, Malcolm Hardy, and Mykel J. Kochenderfer. BetterBench: Assessing AI Benchmarks, Uncovering Issues, and Establishing Best Practices. In *Advances in Neural Information Processing Systems*, volume 37, 2024.
- [20] Amaia Salvador, Nicholas Hynes, Yusuf Aytar, Javier Marin, Ferda Ofli, Ingmar Weber, and Antonio Torralba. Learning Cross-Modal Embeddings for Cooking Recipes and Food Images. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 3020–3028, 2017.
- [21] Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyang Jiang, Tianle Li, Max Ku, Kai Wang, Alex Zhuang, Rongqi Fan, Xiang

- Yue, and Wenhui Chen. MMLU-Pro: A More Robust and Challenging Multi-Task Language Understanding Benchmark. In *Advances in Neural Information Processing Systems*, volume 37, 2024.
- [22] Colin White, Samuel Dooley, Manley Roberts, Arka Pal, Benjamin Feuer, Siddhartha Jain, Ravid Shwartz-Ziv, Neel Jain, Khalid Saifullah, Sreemanti Dey, Shubh Agrawal, Sandeep Singh Sandha, Siddhartha Venkat Naidu, Chinmay Hegde, Yann LeCun, Tom Goldstein, Willie Neiswanger, and Micah Goldblum. LiveBench: A Challenging, Contamination-Limited LLM Benchmark. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [23] Zirui Wu, Xiao Liu, Jiayi Li, Lingpeng Kong, and Yansong Feng. Recipe2Plan: Evaluating Planning Abilities of LLMs for Efficient and Feasible Multitasking with Time Constraints Between Actions. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 4279–4301, Suzhou, China, 2025. Association for Computational Linguistics.
- [24] Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan.  $\tau$ -bench: A Benchmark for Tool-Agent-User Interaction in Real-World Domains. In *The Thirteenth International Conference on Learning Representations*, 2025.