


FlavourBench: executable culinary reasoning without a model judge



THE RESULT

GPT-5.6 Sol Pro

highest score in this run: 62.5

20 frontier models

32 Epicure-generated tasks

640 matched model and Epicure pairs

45.3 to 77.1%

leader Wilson 95% interval

Score only ranks | exact offline replay | direct Kimi and Cohere