
FLAVOURBENCH: An Executable Benchmark for Culinary Reasoning Without a Model Judge

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 FLAVOURBENCH evaluates culinary reasoning against answer keys computed by
2 EPICURE, a versioned executable culinary system. Four deterministic, read-only
3 operations compile a common 32-task panel before any evaluated model is called;
4 no human or model judge scores the test. We run 20 current endpoints on every
5 task in two matched conditions. Model only yields the FlavourBench Score. Model
6 + Epicure permits one read-only call to the specified operation, and the accuracy
7 difference defines Epicure Gain. The release contains 640 matched pairs and
8 1280 assigned arms. GPT-5.6 Sol Pro leads the model-only ranking at 62.5%.
9 Across the panel, Epicure access raises exact accuracy from 43.9% to 93.9%, a
10 +50.0-point gain. Because the operation is named in the prompt, this condition
11 tests tool execution rather than tool discovery. A content-addressed release binds
12 every task, answer key, route, response, completion state, and Epicure trace. Its
13 dependency-free verifier reconstructs the leaderboard and paper figures offline.
14 FLAVOURBENCH measures model knowledge and the value of a domain system
15 on the same task set.

16 1 Introduction

17 Benchmarks are easiest to trust when their ground truth is external to the models being tested. Code
18 can be executed, database states can be inspected, and games can be scored. Culinary reasoning
19 rarely has that luxury. Ingredient substitutions, pairings, technique choices, and cuisine associations
20 are commonly assessed by another language model or by a small set of human preferences. The
21 evaluator then becomes part of the construct.

22 FLAVOURBENCH starts from EPICURE, a read-only culinary runtime with 1,790 ingredients repre-
23 sented in a 300-dimensional space [1, 2]. Four runtime operations are converted into four-choice
24 questions. The compiler records each answer before it contacts an evaluated model.

25 The design asks two separate questions. The FlavourBench Score measures how often a model
26 endpoint returns EPICURE’s answer without access to the runtime. Epicure Gain measures the change
27 when the same endpoint receives the same task with one read-only EPICURE call. Every model
28 receives every task in both conditions. The score determines the leaderboard; Epicure Gain does not
29 enter the ranking.

30 Our release evaluates 20 current endpoints, including GPT-5.6 Sol, Terra, and Luna; Claude Fable,
31 Opus, and Sonnet 5; Gemini 3.1 Pro and 3.6 Flash; Grok 4.5; Kimi K3; Qwen 3.8 Max; GLM 5.2;
32 two DeepSeek V4 variants; MiniMax M3; Nemotron 3 Ultra; Mistral Large 2512; Tencent HY 3;
33 and Cohere Command A Plus and Command A Reasoning. Kimi and both Cohere models use direct




Figure 1: The FLAVOURBENCH pipeline. EPICURE fixes each answer before evaluation. The Model only arm produces the ranked FlavourBench Score. The matched Model + Epicure arm produces Epicure Gain. Both are reconstructed from the same public release.

Family	Epicure operation	Exact target	Example query	Tasks
Substitution	<code>neighbors</code>	first ranked neighbour	top neighbour for tomato	8
Composition	<code>pairing_score</code>	four-decimal affinity	tomato + basil	8
Cookability	<code>compare_on_axis</code>	larger flavour projection	coffee vs. chocolate on bitter	8
Evidence	<code>cultural_profile</code>	highest cuisine cosine	cuisine direction for kimchi	8

Table 1: The task set is balanced by construction. Each family maps one deterministic EPICURE operation to a four-choice answer. Chance accuracy is 25%.

34 provider routes. Every other endpoint uses an exact, no-fallback OpenRouter route selected before
 35 the run.

36 The paper makes three contributions:

- 37 1. an automated culinary benchmark whose questions and answer keys come from a versioned
- 38 executable system rather than a model judge;
- 39 2. a common-task frontier evaluation with 640 matched pairs, exact missing-response penalties, and
- 40 a public 20-row leaderboard;
- 41 3. a matched Epicure intervention that reports panel gain and task-family interactions without
- 42 changing the ranking score.

43 EPICURE has two roles. It fixes the answer key, then becomes the intervention in the second arm.
 44 This keeps the leaderboard and the Epicure comparison connected without combining them into an
 45 arbitrary composite score.

46 2 Epicure reference answers

47 2.1 A versioned reference

48 For FLAVOURBENCH, ground truth is the output of a specific version of EPICURE. It is not a claim
 49 about universal culinary preference. The release binds the application, evidence bundle, release
 50 identifier, ingredient count, and embedding dimensionality by hash. A different runtime produces a
 51 different task-set hash.

52 This definition makes every answer independently recomputable. It also gives the score a direct
 53 interpretation: accuracy is agreement with the published culinary runtime under the published task
 54 construction. No second model decides which response sounds better.

55 2.2 Four balanced task families

56 The compiler creates eight tasks in each of four families. Seeds and ingredient pairs are fixed before
 57 evaluation. For every task, EPICURE is called once to obtain the reference result. The correct answer
 58 is extracted programmatically, three distractors are constructed without model generation, and the
 59 four options are shuffled with a task-ID-derived seed.

60 The task families deliberately mix retrieval and computation. Neighbour and cuisine questions test
 61 whether a model knows relationships encoded by the graph. Pairing questions require an exact
 62 numeric result that is difficult to guess. Axis questions test a comparative claim rather than a

63 memorised scalar. Together they prevent the leaderboard from reducing to one operation or one kind
64 of culinary trivia.

65 2.3 Prompt and answer contract

66 Each prompt names the relevant operation and gives four choices. Models may explain briefly, but
67 the last answer line must be FINAL_CHOICE: X. The parser accepts A, B, C, or D on that marker
68 line. A missing marker, abnormal provider finish, timeout, or absent response scores zero. The rule
69 fits in the standalone verifier shipped with the release.

70 The complete task set contains no hidden human label and no model-authored item. The benchmark is
71 therefore suitable for automatic reruns, CI-style regression tests, and training loops where thousands
72 of candidate responses must be scored consistently.

73 3 Evaluation design

74 3.1 One score and one matched intervention

75 For model m and task t , let M_{mt} equal one when the Model only response finishes normally and
76 matches the EPICURE answer. Let E_{mt} be the same indicator for Model + Epicure. The prompt,
77 route, temperature, top- p , seed request, output limit, and answer parser are identical. Only access to
78 EPICURE changes.

79 Each Model + Epicure response may make one read-only call to the operation named in the prompt.
80 The runner records the operation, arguments, latency, error flag, and a hash of the returned payload.
81 The returned data is bounded before it enters model context. Automatic provider fallback is disabled,
82 so a route failure remains a failure for that route.

83 The single ranking metric is the FlavourBench Score:

$$F_m = 100 \frac{1}{32} \sum_t M_{mt}. \quad (1)$$

84 Every family contains eight tasks, so ordinary accuracy already gives each family equal weight.
85 Missing or invalid responses remain in the denominator as incorrect. The public release sorts by F_m ,
86 then uses Model + Epicure accuracy, Model + Epicure completion, and model ID as fixed tie-breakers.

87 Epicure Gain is the matched difference

$$G_m = 100 \frac{1}{32} \sum_t (E_{mt} - M_{mt}), \quad (2)$$

88 reported in percentage points. G_m is not part of the FlavourBench Score. The release also keeps the
89 four paired outcomes for every task: both correct, Model only correct, Model + Epicure correct, or
90 neither correct.

91 3.2 Completion and uncertainty

92 Completion is the fraction of assigned responses that end normally with a parseable final marker.
93 The score still counts missing responses as incorrect. Reporting completion beside the score keeps
94 endpoint failure separate from a wrong culinary answer. A zero score with zero completions is not
95 evidence of zero model capability. For each 32-task accuracy we report a 95% Wilson interval.

96 The benchmark is designed for visible differences, not hundredths of a point. One answer changes
97 the score by 3.125 points. Family breakdowns and the full pair matrix are published so tied or nearby
98 aggregate scores can be inspected rather than over-interpreted.

99 3.3 Twenty-model frontier panel

100 The panel was fixed before the final calls. It includes three GPT-5.6 tiers rather than treating one
101 deployment as the entire OpenAI family; three Claude 5 tiers; two Gemini tiers; two DeepSeek
102 V4 routes; both current Cohere Command A routes; and a broad set of current independent and




Figure 2: FlavourBench Score against EPICURE answer keys. Points are exact accuracy over the common 32-task panel; bars are 95% Wilson intervals. Crosses mark endpoints with no Model only completion. The dashed line is four-choice chance.

open-weight families. Kimi K3 runs through Kimi’s Anthropic-compatible endpoint. Command A Plus and Command A Reasoning run through Cohere V2 Chat. Qwen 3.8 Max and GLM 5.2 use explicit OpenRouter routes in the automated batch [5–7, 14, 18, 19]. The exact 20-row route table appears in Appendix C.

4 Results

4.1 The FlavourBench Score leaderboard

GPT-5.6 Sol Pro leads with 20/32 correct and a FlavourBench Score of 62.5. Table 2 gives the full release ordering and completion rates. Figure 2 adds intervals to show which gaps the 32-task panel can resolve. Models with no Model only completion are marked DNF rather than interpreted as the least capable systems.

The answer key exists before the model call, every endpoint receives the same questions, and no judge is sampled after the fact. Missing outputs count as incorrect in the operational score but remain visible in the completion columns. Mistral Large 2512 and Nemotron 3 Ultra returned no Model only completion in this run, so their zeroes are transport outcomes rather than capability estimates. Nemotron also returned no Model + Epicure completion.

4.2 What Epicure changes

Across all 640 matched pairs, panel accuracy is 43.9% for Model only and 93.9% for Model + Epicure. Epicure Gain is therefore +50.0 points. Every route that produced at least one assisted response gains points under the operational scoring rule. The largest numerical change is Mistral Large 2512 at

Rank	Model	FB Score	+Epicure	Gain	Model % done	Epicure % done
1	GPT-5.6 Sol Pro	62.5	100.0	+37.5	96.9	100.0
2	Claude Sonnet 5	59.4	100.0	+40.6	100.0	100.0
3	DeepSeek V4 Pro	59.4	93.8	+34.4	100.0	93.8
4	Gemini 3.6 Flash	56.2	100.0	+43.8	100.0	100.0
5	Kimi K3	56.2	100.0	+43.8	100.0	100.0
6	Tencent HY 3	56.2	100.0	+43.8	100.0	100.0
7	Claude Fable 5	53.1	100.0	+46.9	93.8	100.0
8	Claude Opus 5	50.0	100.0	+50.0	100.0	100.0
9	GPT-5.6 Luna Pro	46.9	100.0	+53.1	100.0	100.0
10	GPT-5.6 Terra Pro	46.9	100.0	+53.1	100.0	100.0
11	Qwen 3.8 Max	46.9	100.0	+53.1	87.5	100.0
12	Command A Reasoning	46.9	96.9	+50.0	93.8	100.0
13	Gemini 3.1 Pro	46.9	96.9	+50.0	93.8	96.9
14	MiniMax M3	43.8	100.0	+56.2	96.9	100.0
15	GLM 5.2	43.8	96.9	+53.1	100.0	96.9
16	Command A Plus	40.6	100.0	+59.4	100.0	100.0
17	DeepSeek V4 Flash	40.6	100.0	+59.4	75.0	100.0
18	Grok 4.5	21.9	100.0	+78.1	31.2	100.0
DNF	Mistral Large 2512	0.0	93.8	+93.8	0.0	100.0
DNF	Nemotron 3 Ultra	0.0	0.0	+0.0	0.0	0.0

Table 2: The public leaderboard. “FB Score” is Model only exact-choice accuracy and determines rank. “+Epicure” is accuracy for the same tasks with EPICURE available. “Gain” is their difference in percentage points. The final columns report completion. DNF means that the Model only endpoint returned no parseable completion.

Family	Epicure operation	Tasks	Model only	Model + Epicure
Substitution	neighbors	8	18.1	94.4
Composition	pairing_score	8	40.6	94.4
Cookability	compare_on_axis	8	62.5	92.5
Evidence	cultural_profile	8	54.4	94.4

Table 3: Panel-average performance by family and the corresponding reference operation.

+93.8 points, but that route had no Model only completion. Its change combines condition-specific availability with answer quality and should not be read as a pure knowledge gain.

The assisted prompt names the relevant operation. This is a scaffolded test of calling that operation, reading its output, and preserving the answer contract. It is not a test of open-ended tool discovery. That narrower intervention is useful here because it isolates the benefit of EPICURE from the separate problem of choosing among many tools.

4.3 Families explain aggregate ties

The four task families have different demands. Pairing questions require an exact four-decimal value; substitution and cultural questions can sometimes be answered from broad food knowledge; axis comparisons require a directional decision in EPICURE’s latent space. Figure 4 shows that similar total scores can arise from different profiles.

4.4 Every model and task pair is inspectable

Aggregate bars can hide a benchmark bug or a single pathological task. Figure 5 therefore publishes all 640 paired outcomes. Blue cells are stable knowledge: both conditions are correct. Gold cells are Epicure wins: Model only is wrong and Model + Epicure is correct. Red cells show the reverse outcome. Grey cells show questions neither condition solved.

4.5 Cohere, Kimi, Qwen, and route fidelity

Both Cohere models are full members of the public panel. Command A Plus ranks 16 with a FlavourBench Score of 40.6 and an Epicure Gain of +59.4 points. Command A Reasoning ranks 12




Figure 3: Model only and Model + Epicure accuracy over the same 32 tasks. Each line is a matched comparison within one endpoint, not a comparison between different task samples.

141 with a score of 46.9 and a gain of +50.0 points. Kimi K3 ranks 5 through the direct Kimi endpoint.
 142 Qwen 3.8 Max ranks 11 through an exact OpenRouter route, and GLM 5.2 ranks 15. No route
 143 changes during generation.

144 The two direct Cohere routes each produced all 64 assigned arms. Their returned identifiers are comm
 145 and-a-plus-05-2026 and command-a-reasoning-08-2025, respectively; the direct Kimi route
 146 likewise produced all 64 arms under returned identifier k3. The offline verifier checks these provider,
 147 identity, and arm-count contracts before accepting the release.

148 This matters because a model name is not a transport contract. The release stores the requested model,
 149 returned model, provider, route, generation ID, latency, token counts, and response hash. Direct and
 150 routed endpoints can therefore be analysed separately without pretending they are interchangeable
 151 deployments.

152 4.6 Score, response time, and Epicure Gain

153 Figure 6 plots FlavourBench Score against median response time. Marker colour encodes Epicure
 154 Gain. A fast model can score poorly, a high-scoring model can gain little from EPICURE, and a
 155 lower-ranked model can improve sharply when assisted. The release keeps these dimensions separate.

156 5 Using the score as a training signal

157 The same environment exposes four binary indicators for a generated trajectory τ : r_{answer} for a
 158 correct answer, r_{format} for a parseable final marker, r_{call} for a successful Epicure call, and r_{align} for
 159 a call matching the reference operation. The public leaderboard uses r_{answer} in Model only. A future

160 training objective could use

$$r(\tau) = r_{\text{answer}}r_{\text{format}} + \lambda r_{\text{answer}}r_{\text{call}} + \mu r_{\text{align}}, \quad (3)$$

161 where each component is a binary indicator and λ, μ set the value of successful and reference-
162 matching calls. Requiring a correct answer prevents irrelevant calls from earning the main reward.
163 The interface could support rejection sampling, preference optimisation, or reinforcement learning on
164 new ingredient seeds. This paper does not train a checkpoint; such a study should keep the public
165 leaderboard as a final evaluation split.

166 6 Reproducibility and release

167 The public release is one JSON document with hash 197b84d75cca... It contains the 20 route
168 records, 32 full prompts, answer keys and reference EPICURE results, 1,280 observation rows, final
169 answer markers, response times, and bounded Epicure-trace projections. 1195 provider responses
170 were observed; 85 missing arms remain explicit zero-scored rows.

171 The leaderboard can be reconstructed without a network connection, provider key, GPU, or running
172 EPICURE service:

```
173 python3 -I reproduce_epicure_native.py \  
174 --release generated/epicure-native/  
175 epicure-native-release.json
```

176 The verifier rejects duplicate JSON keys, non-finite numbers, a changed content address, an incom-
177 plete $20 \times 32 \times 2$ grid, or any leaderboard value that differs from recomputation. The paper tables
178 and figures are generated from the same release, so a changed score cannot leave a stale chart behind.

179 Reproducing the original hosted calls is a different operation from replaying the result. The release
180 records exact routes and request policies, but provider deployments can change. A fresh frontier run
181 should generate a new manifest and release rather than overwrite this snapshot.

182 7 Discussion

183 **What the score means.** The FlavourBench Score answers a narrow operational question: how
184 often did this deployed endpoint return the answer fixed by the published EPICURE runtime? It
185 includes endpoint availability because missing responses score zero. The completion columns reveal
186 that component. The score is not a universal measure of cooking skill, but it is exact, reproducible,
187 and independent of a model judge.

188 **Do we need humans?** Not for this leaderboard. Humans are useful when the target is preference,
189 creativity, safety, or the quality of an open-ended recipe. They are unnecessary when the question
190 is whether a model’s choice matches a deterministic runtime result. FLAVOURBENCH uses that
191 distinction to automate the core ranking while leaving subjective culinary evaluation as a separate
192 track.

193 **Why run Model + Epicure?** The leaderboard says which endpoint best predicts EPICURE. The
194 matched intervention asks how much the same endpoint changes when given EPICURE directly.
195 That result is operationally useful: a strong model may gain little, while a lower-scoring model may
196 become competitive when grounded. The dumbbell and pair matrix show this task by task.

197 **Scaling the benchmark.** The task compiler can expand the seed lists while preserving the four
198 operations. The next release should reserve unseen ingredients and pairs, remove near-duplicate
199 questions, and measure stability across runtime versions. The current 32-task panel is a frontier
200 snapshot and a reproducible template for a larger hidden set.

201 8 Related work

202 HELM argues for explicit scenarios and metrics, while recent benchmark audits emphasise validity,
203 contamination, and transparent reporting [15, 23]. Chatbot Arena and WildBench evaluate broad

Benchmark	Scope	Primary target	Paired tool?
Chatbot Arena	general chat	pairwise human preference	No
WildBench	general chat	task-specific checklist	No
BFCL	function calling	executable call match	No
τ -bench	interactive agents	exact goal state	No
DiningBench	multimodal food	dish and nutrition labels	No
FAM-Bench	food-as-medicine	expert suitability label	No
FLAVOURBENCH	culinary reasoning	versioned executable EPICURE operation	Yes

Table 4: FLAVOURBENCH complements broad, tool-use, and food-domain benchmarks. Its distinctive unit is the same endpoint and task measured as Model only and Model + Epicure; “paired tool” denotes that matched intervention rather than ordinary tool availability.

answer quality through preferences or checklists [3, 16]. FLAVOURBENCH instead targets an executable domain model and needs no judge for its primary score.

ToolLLM, BFCL, and τ -bench make tool invocation or environment state observable [20, 21, 28]. FLAVOURBENCH adds a matched Model only condition, so EPICURE becomes an intervention rather than a requirement shared by every response. LiveBench, LiveCodeBench, GPQA, MMLU-Pro, and SWE-bench illustrate complementary strategies for reducing contamination or grounding answers in difficult tasks and executable outcomes [11, 12, 22, 25, 26].

Culinary datasets such as FlavorDB and Recipe1M organise ingredient and recipe evidence [9, 24]. Recent food benchmarks emphasise other constructs: NutriBench estimates carbohydrates from meal descriptions; FoodBench-QA grounds nutrition tasks in databases and ontologies; DiningBench tests multimodal dish perception and nutrition reasoning; and FAM-Bench evaluates condition-aware food-as-medicine decisions [8, 10, 13, 17]. CookingSense and planning-oriented culinary work test commonsense and procedures [4, 27]. FLAVOURBENCH targets a complementary construct: ingredient-space flavour decisions whose targets are compiled from a callable, versioned runtime, with access to that runtime measured as a matched intervention.

9 Limitations

The 32-task panel resolves large score differences better than close ones: one answer moves the score by 3.125 points. The English ingredient seeds cover four EPICURE operations, and hosted routes are a dated snapshot. Public prompts may eventually be memorised. The assisted condition names the operation, so it measures execution with a specified tool rather than open-ended tool selection. Finally, agreement with EPICURE inherits the runtime’s representation choices; it is not a food-safety or cultural-authenticity certification. Larger hidden panels, additional operations, and repeated route snapshots are the immediate next tests.

10 Conclusion

FLAVOURBENCH turns a versioned culinary runtime into an executable benchmark for frontier language models. GPT-5.6 Sol Pro leads the 20-model panel with a FlavourBench Score of 62.5. Giving the same endpoints a specified EPICURE operation raises panel accuracy by +50.0 points. The score, gain, completion rates, prompts, responses, and traces all reconstruct from the released 1,280-arm grid. The result is a simple benchmark with a separate, measured intervention: what models know, and what EPICURE adds.

234 A Real prompts, answers, and tool calls

235 A benchmark paper should show what was actually asked. The four cases below are selected
236 mechanically, one per family. Where possible, the builder chooses a model and task that is wrong in
237 Model only and correct in Model + Epicure. Each case includes the full prompt, both final answers,
238 the actual operation and arguments, and the exact answer key.

239 **Case 1: Substitution with GPT-5.6 Sol Pro. Prompt.** FlavourBench Epicure-native exact task.
240 Use the read-only culinary evidence tools if they are available; otherwise answer from your own
241 knowledge.

242 According to Epicure’s ‘neighbors’ result for ‘tomato’ with top_k=4, which ingredient is ranked first?

243 Choices: A. oregano B. bell_pepper C. olive_oil D. onion

244 In your final answer, explain briefly if useful and finish on its own line with exactly ‘FINAL_CHOICE:
245 X’, replacing X with A, B, C, or D.

246 **Model only.** FINAL_CHOICE: D **Model + Epicure.** FINAL_CHOICE: B

247 **Epicure call.** neighbors({"ingredient": "tomato", "top_k": 4}). The answer key is
248 FINAL_CHOICE: B.

249 **Epicure result.** {"ingredient": "tomato", "neighbors": [{"name": "bell_pepper", "rank":
250 1, "sim": 0.4799}, {"name": "olive_oil", "rank": 2, "sim": 0.4786}, {"name": "oregano",
251 "rank": 3, "sim": 0.4646}, {"name": "onion", "rank": 4, "sim": 0.4567}]}

252 **Case 2: Composition with DeepSeek V4 Pro. Prompt.** FlavourBench Epicure-native exact task.
253 Use the read-only culinary evidence tools if they are available; otherwise answer from your own
254 knowledge.

255 What exact four-decimal ‘pairing_score’ does Epicure return for ‘tomato’ and ‘basil’?

256 Choices: A. 0.2707 B. 0.2407 C. 0.3207 D. 0.1907

257 In your final answer, explain briefly if useful and finish on its own line with exactly ‘FINAL_CHOICE:
258 X’, replacing X with A, B, C, or D.

259 **Model only.** FINAL_CHOICE: B **Model + Epicure.** FINAL_CHOICE: A

260 **Epicure call.** pairing_score({"ingredient_a": "tomato", "ingredient_b":
261 "basil"}). The answer key is FINAL_CHOICE: A.

262 **Epicure result.** {"all_pairs_range": {"median": 0.092, "p10": 0.0197, "p25": 0.0521, "p75":
263 0.1392, "p90": 0.1899}, "pairing_score": 0.2707, "percentile_label": "very high (>=p90)",
264 "resolved_a": "tomato", "resolved_b": "basil"}

265 **Case 3: Cookability with GPT-5.6 Sol Pro. Prompt.** FlavourBench Epicure-native exact task.
266 Use the read-only culinary evidence tools if they are available; otherwise answer from your own
267 knowledge.

268 Which statement matches Epicure’s ‘compare_on_axis’ result for ‘coffee’ versus ‘chocolate’ on
269 ‘cf_bitter’?

270 Choices: A. coffee has the higher projection B. the projections are equal to four decimals C. Epicure
271 cannot resolve one or both ingredients D. chocolate has the higher projection

272 In your final answer, explain briefly if useful and finish on its own line with exactly ‘FINAL_CHOICE:
273 X’, replacing X with A, B, C, or D.

274 **Model only.** FINAL_CHOICE: A **Model + Epicure.** FINAL_CHOICE: D

275 **Epicure call.** compare_on_axis({"axis": "cf_bitter", "ingredient_a": "coffee",
276 "ingredient_b": "chocolate"}). The answer key is FINAL_CHOICE: D.

277 **Epicure result.** {"axis": "cf_bitter", "axis_range": {"p10": -0.0877, "p90": 0.1429},
278 "delta": 0.0835, "label_a": "moderate", "label_b": "high", "projection_a": 0.0076,
279 "projection_b": 0.0911, "resolved_a": "coffee", "resolved_b": "chocolate"}

280 **Case 4: Evidence with Grok 4.5. Prompt.** FlavourBench Epicure-native exact task. Use the
281 read-only culinary evidence tools if they are available; otherwise answer from your own knowledge.
282 Which cuisine direction has the highest cosine score in Epicure’s ‘cultural_profile’ result for ‘miso’?
283 Choices: A. Japanese B. East_Asian C. Southeast_Asian D. South_Asian
284 In your final answer, explain briefly if useful and finish on its own line with exactly ‘FINAL_CHOICE:
285 X’, replacing X with A, B, C, or D.
286 **Model only.** {"name": "cultural_profile", "arguments": {"ingredient":
287 "miso"}} **Model + Epicure.** FINAL_CHOICE: A
288 **Epicure call.** cultural_profile({"ingredient": "miso"}). The answer key is
289 FINAL_CHOICE: A.
290 **Epicure result.** {"cuisines": {"East_Asian": {"range": {"p10": -0.1649, "p90":
291 0.3039}, "score": 0.1625}, "Eastern_European": {"range": {"p10": -0.1466, "p90":
292 0.1034}, "score": -0.1494}, "Japanese": {"range": {"p10": -0.1359, "p90": 0.1896},
293 "score": 0.3227}, "Latin_American": {"range": {"p10": -0.1672, "p90": 0.1281},
294 "score": -0.1538}, "Mediterranean": {"range": {"p10": -0.2323, "p90": 0.206}, "score":
295 -0.1198}, "South_Asian": {"range": {"p10": -0.0808, "p90": 0.071}, "score": -0.0585},
296 "Southeast_Asian": {"range": {"p10": -0.1625, "p90": 0.193}, "score": 0.0426},
297 "Western_Atlantic": {"range": {"p10": -0.2304, "p90": 0.183}, "score": -0.1546}},
298 "note": "Cosine similarity to each cuisine direction. Use p10/p90 to judge whether a score
299 is typical or extreme.", "resolved": "miso"}
300 These examples show exactly what the two measurements contain. Model only tests knowledge or
301 inference about the hidden graph. Model + Epicure names the operation, so it tests whether the model
302 can form valid arguments, read the returned structure, and preserve the answer contract. A correct
303 choice with the wrong operation still earns answer credit, but the release separately reports whether
304 the reference operation was matched. Across the run there are 606 projected EPICURE calls and 606
305 model and task pairs matching the reference operation.

306 B Additional result figures

307 C Exact execution routes

308 D Protocol details

309 All arms request temperature 0, top-*p* 1, and seed 20260810 where the provider supports it. The final
310 response budget is 2,048 tokens. Model + Epicure has one call round, one call total, a 4,096-byte
311 result bound, and a 4,096-byte cumulative result bound. At most two provider attempts are allowed,
312 and an ambiguously delivered request is never replayed. The final marker is parsed from plain text for
313 transport portability; direct providers may use native structured generation internally and normalize
314 to the same marker.

315 E Data fields

316 Each observation binds model ID, task ID, condition, source status, source and response hashes,
317 returned model and provider, finish reason, response time, final answer, observed and expected choice,
318 normal-parse flag, correctness, and a projected Epicure trace. Each trace event contains its name,
319 structured arguments, error flag, response time, and result hash. The release does not require raw
320 provider credentials or hidden service state.




Figure 4: Model only and Model + Epicure accuracy by task family. Each value is a percentage over eight tasks. Similar FlavourBench Scores can come from different family profiles.




Figure 5: All 640 matched outcomes, ordered by FlavourBench Score and grouped by task family. The matrix shows which exact tasks produce Epicure Gain and where the assisted response regresses.




Figure 6: FlavourBench Score, response time, and Epicure Gain. Response time uses a log scale. Point size increases with the magnitude of gain; colour indicates its direction. An × marks a route with no observed response time.

Slot	Model	Execution route
1	GPT-5.6 Sol Pro	openrouter: openai/flex
6	Claude Sonnet 5	openrouter: amazon-bedrock/claude-on-aws
13	DeepSeek V4 Pro	openrouter: baidu/fp8
8	Gemini 3.6 Flash	openrouter: google-vertex/global/flex
10	Kimi K3	kimi_direct: kimi-code-direct
18	Tencent HY 3	openrouter: baidu/fp8
4	Claude Fable 5	openrouter: amazon-bedrock
5	Claude Opus 5	openrouter: amazon-bedrock/claude-on-aws
3	GPT-5.6 Luna Pro	openrouter: openai/flex
2	GPT-5.6 Terra Pro	openrouter: openai/flex
11	Qwen 3.8 Max	openrouter: alibaba
20	Command A Reasoning	cohere_direct: cohere-direct
7	Gemini 3.1 Pro	openrouter: google-vertex/global/flex
15	MiniMax M3	openrouter: together
12	GLM 5.2	openrouter: coreweave/fp4
19	Command A Plus	cohere_direct: cohere-direct
14	DeepSeek V4 Flash	openrouter: cloudflare/fp8
9	Grok 4.5	openrouter: xai/zdr
17	Mistral Large 2512	openrouter: mistral
16	Nemotron 3 Ultra	openrouter: deepinfra/fp4

Table 5: All 20 model routes in the released leaderboard. Routes are selected before generation; automatic fallback is disabled.

References

- [1] Anonymous. Epicure: Navigating the Emergent Geometry of Food Ingredient Embeddings. *arXiv preprint arXiv:2605.22391*, 2026.
- [2] Anonymous. Epicure: Multidimensional Flavor Structure in Food Ingredient Embeddings. *arXiv preprint arXiv:2604.22776*, 2026.
- [3] Wei-Lin Chiang, Lianmin Zheng, Ying Sheng, Anastasios Nikolas Angelopoulos, Tianle Li, Dacheng Li, Hao Zhang, Banghua Zhu, Michael I. Jordan, Joseph E. Gonzalez, and Ion Stoica. Chatbot Arena: An Open Platform for Evaluating LLMs by Human Preference. In *Proceedings of the 41st International Conference on Machine Learning*, volume 235 of *Proceedings of Machine Learning Research*, pages 8359–8388. PMLR, 2024. URL <https://proceedings.mlr.press/v235/chiang24b.html>.
- [4] Donghee Choi, Mogan Gim, Donghyeon Park, Mujeen Sung, Hyunjae Kim, Jaewoo Kang, and Jihun Choi. CookingSense: A Culinary Knowledgebase with Multidisciplinary Assertions. In *Proceedings of the 2024 Joint International Conference on Computational Linguistics, Language Resources and Evaluation (LREC-COLING 2024)*, pages 3983–3996, Torino, Italia, 2024. ELRA and ICCL. URL <https://aclanthology.org/2024.lrec-main.354/>.
- [5] Cohere. Chat (V2) API Reference. <https://docs.cohere.com/v2/reference/chat>, 2026. URL <https://docs.cohere.com/v2/reference/chat>. Accessed 3 August 2026.
- [6] Cohere. Reasoning models. <https://docs.cohere.com/v2/docs/reasoning>, 2026. URL <https://docs.cohere.com/v2/docs/reasoning>. Accessed 3 August 2026.
- [7] Cohere. Tool use overview. <https://docs.cohere.com/v2/docs/tool-use-overview/>, 2026. URL <https://docs.cohere.com/v2/docs/tool-use-overview/>. Accessed 3 August 2026.
- [8] Tome Eftimov, Ana Gjorgjevikj, Matej Martinc, Gjorgjina Cenikj, Saso Dzeroski, and Barbara Koroušić Seljak. FoodBench-QA: Overview of the Shared Task on Grounded Food and Nutrition Question Answering. In *Proceedings of the Third Workshop on Patient-Oriented Language Processing (CL4Health) at LREC 2026*, 2026. URL <https://cs.ijs.si/publication/751>.
- [9] Neelansh Garg, Apuroop Sethupathy, Rudraksh Tuwani, et al. FlavorDB: A Database of Flavor Molecules. *Nucleic Acids Research*, 46(D1):D1210–D1216, 2018. doi: 10.1093/nar/gkx957.
- [10] Andong Hua, Mehak Preet Dhaliwal, Laya Pullela, Ryan Burke, and Yao Qin. NutriBench: A Dataset for Evaluating Large Language Models on Nutrition Estimation from Meal Descriptions. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://arxiv.org/abs/2407.12843>.
- [11] Naman Jain, King Han, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar-Lezama, Koushik Sen, and Ion Stoica. LiveCodeBench: Holistic and Contamination Free Evaluation of Large Language Models for Code. *arXiv preprint arXiv:2403.07974*, 2024. URL <https://arxiv.org/abs/2403.07974>.
- [12] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can Language Models Resolve Real-World GitHub Issues? In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=VTF8yNQm66>.
- [13] Song Jin, Juntian Zhang, Xun Zhang, Zeying Tian, Fei Jiang, Guojun Yin, Wei Lin, Yong Liu, and Rui Yan. DiningBench: A Hierarchical Multi-view Benchmark for Perception and Reasoning in the Dietary Domain. In *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 35289–35310. Association for Computational Linguistics, 2026. doi: 10.18653/v1/2026.acl-long.1630. URL <https://aclanthology.org/2026.acl-long.1630/>.
- [14] Kimi Team. Kimi K3: Open Frontier Intelligence. *arXiv:2607.24653*, 2026. URL <https://arxiv.org/abs/2607.24653>.
- [15] Percy Liang, Rishi Bommasani, Tony Lee, et al. Holistic Evaluation of Language Models. *Transactions on Machine Learning Research*, 2023.
- [16] Bill Yuchen Lin, Yuntian Deng, Khyathi Chandu, Faeze Brahman, Abhilasha Ravichander, Valentina Pyatkin, Nouha Dziri, Ronan Le Bras, and Yejin Choi. WildBench: Benchmarking LLMs with Challenging Tasks from Real Users in the Wild. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/771155abaae744e08576f1f3b4b7ac0d-Abstract-Conference.html.

- [17] Mingyang Mao, Bhargav Rishi Mediseti, Utkarsh Grover, Tanvir Ibrahim, Wenyan Li, Tingting Zhang, and Xiaomin Lin. FAM-Bench: A Multimodal Benchmark for Condition-Aware Food-as-Medicine Reasoning. *arXiv preprint arXiv:2605.31410*, 2026. URL <https://arxiv.org/abs/2605.31410>.
- [18] Moonshot AI. Kimi code: Providers and models. <https://moonshotai.github.io/kimi-code/en/configuration/providers.html>, 2026. URL <https://moonshotai.github.io/kimi-code/en/configuration/providers.html>. Accessed 28 July 2026.
- [19] OpenRouter. Get Request and Usage Metadata for a Generation. <https://openrouter.ai/docs/api/reference/overview>, 2026. URL <https://openrouter.ai/docs/api/reference/overview>. Accessed 2026-07-15.
- [20] Shishir G. Patil, Huanzhi Mao, Fanjia Yan, Charlie Cheng-Jie Ji, Vishnu Suresh, Ion Stoica, and Joseph E. Gonzalez. The Berkeley Function Calling Leaderboard (BFCL): From Tool Use to Agentic Evaluation of Large Language Models. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, pages 48371–48392. PMLR, 2025. URL <https://proceedings.mlr.press/v267/patil25a.html>.
- [21] Yujia Qin, Shihao Liang, Yining Ye, et al. ToolLLM: Facilitating Large Language Models to Master 16000+ Real-world APIs. In *The Twelfth International Conference on Learning Representations*, 2024. URL <https://openreview.net/forum?id=dHng200Jjr>.
- [22] David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R. Bowman. GPQA: A Graduate-Level Google-Proof Q&A Benchmark. In *First Conference on Language Modeling*, 2024. URL <https://openreview.net/forum?id=Ti67584b98>.
- [23] Anka Reuel, Amelia Hardy, Chandler Smith, Max Lamparth, Malcolm Hardy, and Mykel J. Kochenderfer. BetterBench: Assessing AI Benchmarks, Uncovering Issues, and Establishing Best Practices. In *Advances in Neural Information Processing Systems*, volume 37, 2024. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/26889e8359e7ef8a7f5d77457364ca55-Abstract-Datasets_and_Benchmarks_Track.html.
- [24] Amaia Salvador, Nicholas Hynes, Yusuf Aytar, Javier Marin, Ferda Ofli, Ingmar Weber, and Antonio Torralba. Learning Cross-Modal Embeddings for Cooking Recipes and Food Images. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 3020–3028, 2017.
- [25] Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyang Jiang, Tianle Li, Max Ku, Kai Wang, Alex Zhuang, Rongqi Fan, Xiang Yue, and Wenhui Chen. MMLU-Pro: A More Robust and Challenging Multi-Task Language Understanding Benchmark. In *Advances in Neural Information Processing Systems*, volume 37, 2024. doi: 10.52202/079017-3018. URL https://proceedings.neurips.cc/paper_files/paper/2024/hash/ad236edc564f3e3156e1b2feafb99a24-Abstract-Datasets_and_Benchmarks_Track.html.
- [26] Colin White, Samuel Dooley, Manley Roberts, Arka Pal, Benjamin Feuer, Siddhartha Jain, Ravid Shwartz-Ziv, Neel Jain, Khalid Saifullah, Sreemanti Dey, Shubh Agrawal, Sandeep Singh Sandha, Siddhartha Venkat Naidu, Chinmay Hegde, Yann LeCun, Tom Goldstein, Willie Neiswanger, and Micah Goldblum. LiveBench: A Challenging, Contamination-Limited LLM Benchmark. In *The Thirteenth International Conference on Learning Representations*, 2025.
- [27] Zirui Wu, Xiao Liu, Jiayi Li, Lingpeng Kong, and Yansong Feng. Recipe2Plan: Evaluating Planning Abilities of LLMs for Efficient and Feasible Multitasking with Time Constraints Between Actions. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 4279–4301, Suzhou, China, 2025. Association for Computational Linguistics. doi: 10.18653/v1/2025.findings-emnlp.229. URL <https://aclanthology.org/2025.findings-emnlp.229/>.
- [28] Shunyu Yao, Noah Shinn, Pedram Razavi, and Karthik Narasimhan. τ -bench: A Benchmark for Tool-Agent-User Interaction in Real-World Domains. In *The Thirteenth International Conference on Learning Representations*, 2025. URL https://proceedings.iclr.cc/paper_files/paper/2025/hash/1b126cc38b8638e07bef37e7b2bb72bf-Abstract-Conference.html.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: Sections 2–7 state the benchmark contract, observed evidence, scoring rule, results, and scope.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 7 discusses construct, task coverage, route stability, reproducibility, and the intended interpretation of the score.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A]

Justification: The paper introduces no theorem or formal theoretical result; the Wilson formula is a standard descriptive confidence interval.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.

480 • The proofs can either appear in the main paper or the supplemental material, but if they appear
481 in the supplemental material, the authors are encouraged to provide a short proof sketch to
482 provide intuition.

483 • Inversely, any informal proof provided in the core of the paper should be complemented by
484 formal proofs provided in appendix or supplemental material.

485 • Theorems and Lemmas that the proof relies upon should be properly referenced.

486 **4. Experimental result reproducibility**

487 Question: Does the paper fully disclose all the information needed to reproduce the main experi-
488 mental results of the paper to the extent that it affects the main claims and/or conclusions of the
489 paper (regardless of whether the code and data are provided or not)?

490 Answer: [Yes]

491 Justification: Sections 3–6 specify the unit, success contract, ranking rule, exact inputs, and
492 verification command for the central result.

493 Guidelines:

494 • The answer [N/A] means that the paper does not include experiments.

495 • If the paper includes experiments, a [No] answer to this question will not be perceived well by
496 the reviewers: Making the paper reproducible is important, regardless of whether the code and
497 data are provided or not.

498 • If the contribution is a dataset and/or model, the authors should describe the steps taken to make
499 their results reproducible or verifiable.

500 • Depending on the contribution, reproducibility can be accomplished in various ways. For
501 example, if the contribution is a novel architecture, describing the architecture fully might
502 suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary
503 to either make it possible for others to replicate the model with the same dataset, or provide
504 access to the model. In general, releasing code and data is often one good way to accomplish
505 this, but reproducibility can also be provided via detailed instructions for how to replicate the
506 results, access to a hosted model (e.g., in the case of a large language model), releasing of a
507 model checkpoint, or other means that are appropriate to the research performed.

508 • While NeurIPS does not require releasing code, the conference does require all submissions
509 to provide some reasonable avenue for reproducibility, which may depend on the nature of the
510 contribution. For example

511 (a) If the contribution is primarily a new algorithm, the paper should make it clear how to
512 reproduce that algorithm.

513 (b) If the contribution is primarily a new model architecture, the paper should describe the
514 architecture clearly and fully.

515 (c) If the contribution is a new model (e.g., a large language model), then there should either
516 be a way to access this model for reproducing the results or a way to reproduce the model
517 (e.g., with an open-source dataset or instructions for how to construct the dataset).

518 (d) We recognize that reproducibility may be tricky in some cases, in which case authors are
519 welcome to describe the particular way they provide for reproducibility. In the case of
520 closed-source models, it may be that access to the model is limited in some way (e.g.,
521 to registered users), but it should be possible for other researchers to have some path to
522 reproducing or verifying the results.

523 **5. Open access to data and code**

524 Question: Does the paper provide open access to the data and code, with sufficient instructions to
525 faithfully reproduce the main experimental results, as described in supplemental material?

526 Answer: [Yes]

527 Justification: The anonymous supplement includes the 32 generated prompts, complete 1,280-arm
528 release, scoring code, vector tables and figures, and a standard-library verifier that reconstructs
529 every reported rank offline.

530 Guidelines:

531 • The answer [N/A] means that paper does not include experiments requiring code.

532 • Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.

533 • While we encourage the release of code and data, we understand that this might not be possible,
534 so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless
535 this is central to the contribution (e.g., for a new open-source benchmark).

536

537 • The instructions should contain the exact command and environment needed to run to reproduce
538 the results. See the NeurIPS code and data submission guidelines ([https://neurips.cc/p](https://neurips.cc/public/guides/CodeSubmissionPolicy)
539 [ublic/guides/CodeSubmissionPolicy](https://neurips.cc/public/guides/CodeSubmissionPolicy)) for more details.

540 • The authors should provide instructions on data access and preparation, including how to access
541 the raw data, preprocessed data, intermediate data, and generated data, etc.

542 • The authors should provide scripts to reproduce all experimental results for the new proposed
543 method and baselines. If only a subset of experiments are reproducible, they should state which
544 ones are omitted from the script and why.

545 • At submission time, to preserve anonymity, the authors should release anonymized versions (if
546 applicable).

547 • Providing as much information as possible in supplemental material (appended to the paper) is
548 recommended, but including URLs to data and code is permitted.

549 **6. Experimental setting/details**

550 Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters,
551 how they were chosen, type of optimizer) necessary to understand the results?

552 Answer: [\[Yes\]](#)

553 Justification: Section 3 and Appendix C specify tasks, routes, paired arms, execution limits, failure
554 handling, identities, and all ranked denominators.

555 Guidelines:

556 • The answer [\[N/A\]](#) means that the paper does not include experiments.

557 • The experimental setting should be presented in the core of the paper to a level of detail that is
558 necessary to appreciate the results and make sense of them.

559 • The full details can be provided either with the code, in appendix, or as supplemental material.

560 **7. Experiment statistical significance**

561 Question: Does the paper report error bars suitably and correctly defined or other appropriate
562 information about the statistical significance of the experiments?

563 Answer: [\[Yes\]](#)

564 Justification: Section 3 defines the two-sided 95% Wilson intervals and the table and Figure 2
565 report exact counts and intervals.

566 Guidelines:

567 • The answer [\[N/A\]](#) means that the paper does not include experiments.

568 • The authors should answer [\[Yes\]](#) if the results are accompanied by error bars, confidence
569 intervals, or statistical significance tests, at least for the experiments that support the main claims
570 of the paper.

571 • The factors of variability that the error bars are capturing should be clearly stated (for example,
572 train/test split, initialization, random drawing of some parameter, or overall run with given
573 experimental conditions).

574 • The method for calculating the error bars should be explained (closed form formula, call to a
575 library function, bootstrap, etc.)

576 • The assumptions made should be given (e.g., Normally distributed errors).

577 • It should be clear whether the error bar is the standard deviation or the standard error of the
578 mean.

579 • It is OK to report 1-sigma error bars, but one should state it. The authors should preferably
580 report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of
581 errors is not verified.

582 • For asymmetric distributions, the authors should be careful not to show in tables or figures
583 symmetric error bars that would yield results that are out of range (e.g., negative error rates).

584 • If error bars are reported in tables or plots, the authors should explain in the text how they were
585 calculated and reference the corresponding figures or tables in the text.

586 **8. Experiments compute resources**

587 Question: For each experiment, does the paper provide sufficient information on the computer
588 resources (type of compute workers, memory, time of execution) needed to reproduce the experi-
589 ments?

590 Answer: [\[No\]](#)

591 Justification: Section 6 records the local verification environment, but hosted model providers
592 did not disclose inference hardware or energy use, preventing complete provider-side compute
593 accounting.

594 Guidelines:

595 • The answer [\[N/A\]](#) means that the paper does not include experiments.

596 • The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud
597 provider, including relevant memory and storage.
598 • The paper should provide the amount of compute required for each of the individual experimental
599 runs as well as estimate the total compute.
600 • The paper should disclose whether the full research project required more compute than the
601 experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it
602 into the paper).

603 **9. Code of ethics**
604 Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS
605 Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?
606 Answer: [Yes]
607 Justification: The work uses paid provider endpoints and a generated task set without recruiting
608 participants; Sections 7 and Limitations disclose safety, route stability, and interpretation risks.
609 Guidelines:
610 • The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
611 • If the authors answer [No], they should explain the special circumstances that require a deviation
612 from the Code of Ethics.
613 • The authors should make sure to preserve anonymity (e.g., if there is a special consideration due
614 to laws or regulations in their jurisdiction).

615 **10. Broader impacts**
616 Question: Does the paper discuss both potential positive societal impacts and negative societal
617 impacts of the work performed?
618 Answer: [Yes]
619 Justification: The Discussion and Limitations sections cover missing-response scoring, leaderboard
620 interpretation, task-set gaming, provider instability, and food-safety misuse.
621 Guidelines:
622 • The answer [N/A] means that there is no societal impact of the work performed.
623 • If the authors answer [N/A] or [No], they should explain why their work has no societal impact
624 or why the paper does not address societal impact.
625 • Examples of negative societal impacts include potential malicious or unintended uses (e.g.,
626 disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deploy-
627 ment of technologies that could make decisions that unfairly impact specific groups), privacy
628 considerations, and security considerations.
629 • The conference expects that many papers will be foundational research and not tied to par-
630 ticular applications, let alone deployments. However, if there is a direct path to any negative
631 applications, the authors should point it out. For example, it is legitimate to point out that
632 an improvement in the quality of generative models could be used to generate Deepfakes for
633 disinformation. On the other hand, it is not needed to point out that a generic algorithm for
634 optimizing neural networks could enable people to train models that generate Deepfakes faster.
635 • The authors should consider possible harms that could arise when the technology is being used
636 as intended and functioning correctly, harms that could arise when the technology is being used
637 as intended but gives incorrect results, and harms following from (intentional or unintentional)
638 misuse of the technology.
639 • If there are negative societal impacts, the authors could also discuss possible mitigation strategies
640 (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for
641 monitoring misuse, mechanisms to monitor how a system learns from feedback over time,
642 improving the efficiency and accessibility of ML).

643 **11. Safeguards**
644 Question: Does the paper describe safeguards that have been put in place for responsible release
645 of data or models that have a high risk for misuse (e.g., pre-trained language models, image
646 generators, or scraped datasets)?
647 Answer: [N/A]
648 Justification: The release contains no trained model, scraped-media dataset, or executable capabil-
649 ity intended for harmful use; it contains bounded exact-choice responses and tool projections.
650 Guidelines:
651 • The answer [N/A] means that the paper poses no such risks.
652 • Released models that have a high risk for misuse or dual-use should be released with necessary
653 safeguards to allow for controlled use of the model, for example by requiring that users adhere
654 to usage guidelines or restrictions to access the model or implementing safety filters.

655 • Datasets that have been scraped from the Internet could pose safety risks. The authors should
656 describe how they avoided releasing unsafe images.

657 • We recognize that providing effective safeguards is challenging, and many papers do not require
658 this, but we encourage authors to take this into account and make a best faith effort.

659 **12. Licenses for existing assets**

660 Question: Are the creators or original owners of assets (e.g., code, data, models), used in the
661 paper, properly credited and are the license and terms of use explicitly mentioned and properly
662 respected?

663 Answer: [Yes]

664 Justification: The manuscript credits providers and cited benchmark/software sources and docu-
665 ments its own CC BY 4.0 and third-party rights boundaries. All 32 prompts are generated by the
666 released task compiler; raw provider credentials and unbounded tool payloads are omitted.

667 Guidelines:

668 • The answer [N/A] means that the paper does not use existing assets.
669 • The authors should cite the original paper that produced the code package or dataset.
670 • The authors should state which version of the asset is used and, if possible, include a URL.
671 • The name of the license (e.g., CC-BY 4.0) should be included for each asset.
672 • For scraped data from a particular source (e.g., website), the copyright and terms of service of
673 that source should be provided.
674 • If assets are released, the license, copyright information, and terms of use in the package should
675 be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for
676 some datasets. Their licensing guide can help determine the license of a dataset.
677 • For existing datasets that are re-packaged, both the original license and the license of the derived
678 asset (if it has changed) should be provided.
679 • If this information is not available online, the authors are encouraged to reach out to the asset’s
680 creators.

681 **13. New assets**

682 Question: Are new assets introduced in the paper well documented and is the documentation
683 provided alongside the assets?

684 Answer: [Yes]

685 Justification: Section 6 documents the aggregate, leaderboard, verifier, result hashes, and exact
686 reproduction command.

687 Guidelines:

688 • The answer [N/A] means that the paper does not release new assets.
689 • Researchers should communicate the details of the dataset/code/model as part of their sub-
690 missions via structured templates. This includes details about training, license, limitations,
691 etc.
692 • The paper should discuss whether and how consent was obtained from people whose asset is
693 used.
694 • At submission time, remember to anonymize your assets (if applicable). You can either create
695 an anonymized URL or include an anonymized zip file.

696 **14. Crowdsourcing and research with human subjects**

697 Question: For crowdsourcing experiments and research with human subjects, does the paper
698 include the full text of instructions given to participants and screenshots, if applicable, as well as
699 details about compensation (if any)?

700 Answer: [N/A]

701 Justification: No crowdsourcing or human-subject study was conducted. The tasks and answer
702 keys are generated deterministically from the versioned Epicure runtime.

703 Guidelines:

704 • The answer [N/A] means that the paper does not involve crowdsourcing nor research with
705 human subjects.
706 • Including this information in the supplemental material is fine, but if the main contribution of
707 the paper involves human subjects, then as much detail as possible should be included in the
708 main paper.
709 • According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other
710 labor should be paid at least the minimum wage in the country of the data collector.

711 **15. Institutional review board (IRB) approvals or equivalent for research with human subjects**

712 Question: Does the paper describe potential risks incurred by study participants, whether such
713 risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals

714 (or an equivalent approval/review based on the requirements of your country or institution) were
 715 obtained?
 716 Answer: [N/A]
 717 Justification: The reported automated study recruited no human participants and collected no
 718 human judgments; the reported score is computed entirely from executable traces and responses.
 719 Guidelines:
 720 • The answer [N/A] means that the paper does not involve crowdsourcing nor research with
 721 human subjects.
 722 • Depending on the country in which research is conducted, IRB approval (or equivalent) may be
 723 required for any human subjects research. If you obtained IRB approval, you should clearly
 724 state this in the paper.
 725 • We recognize that the procedures for this may vary significantly between institutions and
 726 locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for
 727 their institution.
 728 • For initial submissions, do not include any information that would break anonymity (if applica-
 729 ble), such as the institution conducting the review.

730 **16. Declaration of LLM usage**
 731 Question: Does the paper describe the usage of LLMs if it is an important, original, or non-
 732 standard component of the core methods in this research? Note that if the LLM is used only for
 733 writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor,
 734 or originality of the research, declaration is not required.
 735 Answer: [Yes]
 736 Justification: The evaluated LLM endpoints are the experimental subjects. Section 3 and Ap-
 737 pendix C describe their paired role, tool exposure, routing, and failure handling.
 738 Guidelines:
 739 • The answer [N/A] means that the core method development in this research does not involve
 740 LLMs as any important, original, or non-standard components.
 741 • Please refer to our LLM policy in the NeurIPS handbook for what should or should not be
 742 described.